

Lesson 15: Ensemble Methods - Boosting for Regression

Topics Covered:

- 1 Understanding Ensemble Methods: Bagging vs. Boosting
 - 2 Gradient Boosting for Regression (GBR)
 - 3 AdaBoost for Regression
 - 4 Hyperparameter Tuning for Boosting Algorithms
 - 5 Performance Evaluation with MSE, RMSE, MAE, R²
-

1 Bagging vs. Boosting: What's the Difference?

Bagging (Bootstrap Aggregating) and Boosting are both ensemble learning techniques, but they work differently:

Feature	Bagging	Boosting
Purpose	Reduce variance (overfitting)	Reduce bias (underfitting)
Model Dependency	Models trained independently	Models trained sequentially
Weight Adjustment	Equal weight for all models	Adjusts weights based on errors
Example Algorithms	Random Forest	AdaBoost, Gradient Boosting, XGBoost

- ◆ Bagging reduces variance by averaging multiple models.
 - ◆ Boosting reduces bias by training weak models sequentially, each improving upon the previous one.
-

2 Gradient Boosting for Regression (GBR)

Gradient Boosting is an advanced boosting technique where:

- ✓ Each new model **corrects the errors** of the previous model.
- ✓ Uses **gradient descent** to minimize the loss function.
- ✓ Works well with **small and medium datasets**.

❖ Steps in Gradient Boosting

1. Start with a weak model (usually a Decision Tree ♣).
2. Compute the **error residuals** (difference between actual and predicted values).
3. Train a new model to predict these **residuals**.

-
4. Update the final prediction by adding the new model's prediction.
 5. Repeat the process for multiple iterations.
-

3 □ AdaBoost for Regression

AdaBoost (**Adaptive Boosting**) is another boosting method where:

- ✓ It **assigns higher weights** to misclassified points.
 - ✓ Each model is trained sequentially, **focusing more on difficult samples**.
 - ✓ Works well when **data has noise or outliers**.
-

4 □ Implementing Gradient Boosting & AdaBoost for Regression

Let's use **Gradient Boosting** and **AdaBoost** on a **car price dataset**.

❖ Step 1: Import Libraries

```
import pandas as pd
import numpy as np
import time
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import GradientBoostingRegressor, AdaBoostRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Load dataset
df = pd.read_csv('car_price.csv') # Replace with actual dataset

# Select features and target variable
X = df[['mileage', 'year', 'engine_size']] # Example features
y = df['price'] # Target variable

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

❖ Step 2: Train Gradient Boosting Model

```
# Initialize Gradient Boosting model
gbr = GradientBoostingRegressor(n_estimators=500, learning_rate=0.05,
max_depth=4, random_state=42)

# Train the model
gbr.fit(X_train, y_train)

# Predictions
y_train_pred_gbr = gbr.predict(X_train)
```

```
y_test_pred_gbr = gbr.predict(X_test)
```

❖ Step 3: Train AdaBoost Model

```
# Initialize AdaBoost model
ada = AdaBoostRegressor(n_estimators=500, learning_rate=0.05, random_state=42)

# Train the model
ada.fit(X_train, y_train)

# Predictions
y_train_pred_ada = ada.predict(X_train)
y_test_pred_ada = ada.predict(X_test)
```

❖ Step 4: Evaluate Performance (MSE, RMSE, MAE, R²)

```
def evaluate_model(y_true, y_pred, model_name):
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(y_true, y_pred)
    r2 = r2_score(y_true, y_pred)

    print(f"{model_name} Performance:")
    print(f"  MSE: {mse:.2f}")
    print(f"  RMSE: {rmse:.2f}")
    print(f"  MAE: {mae:.2f}")
    print(f"  R2: {r2:.2f}")
    print("-" * 40)

# Evaluate Gradient Boosting
evaluate_model(y_test, y_test_pred_gbr, "Gradient Boosting")

# Evaluate AdaBoost
evaluate_model(y_test, y_test_pred_ada, "AdaBoost")
```

5□ Hyperparameter Optimization (Grid Search)

To further improve performance, we can **tune hyperparameters**.

```
# Define hyperparameters to search for Gradient Boosting
param_grid = {
    'n_estimators': [100, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 4, 5]
}

grid_search = GridSearchCV(GradientBoostingRegressor(), param_grid, cv=3,
                           scoring='r2', n_jobs=-1)
grid_search.fit(X_train, y_train)

print("Best Parameters:", grid_search.best_params_)
```

```
# Train model with best parameters
best_gbr = grid_search.best_estimator_
y_test_pred_best_gbr = best_gbr.predict(X_test)

# Evaluate Optimized Model
evaluate_model(y_test, y_test_pred_best_gbr, "Optimized Gradient Boosting")
```

➤ Key Takeaways

- ✓ **Bagging vs. Boosting:** Bagging reduces variance; Boosting reduces bias.
- ✓ **Gradient Boosting:** Uses gradient descent to minimize errors.
- ✓ **AdaBoost:** Focuses on difficult data points by re-weighting them.
- ✓ **Hyperparameter Tuning:** Improves model performance by finding the best parameters.
- ✓ **Evaluation Metrics:** Always use **MSE**, **RMSE**, **MAE**, **R²** to assess model accuracy.